AGENTIC AI SYSTEMS · HARNESS, LECTURE 2

Advanced Harness Architectures and Their Evaluation

Yonathan Efroni

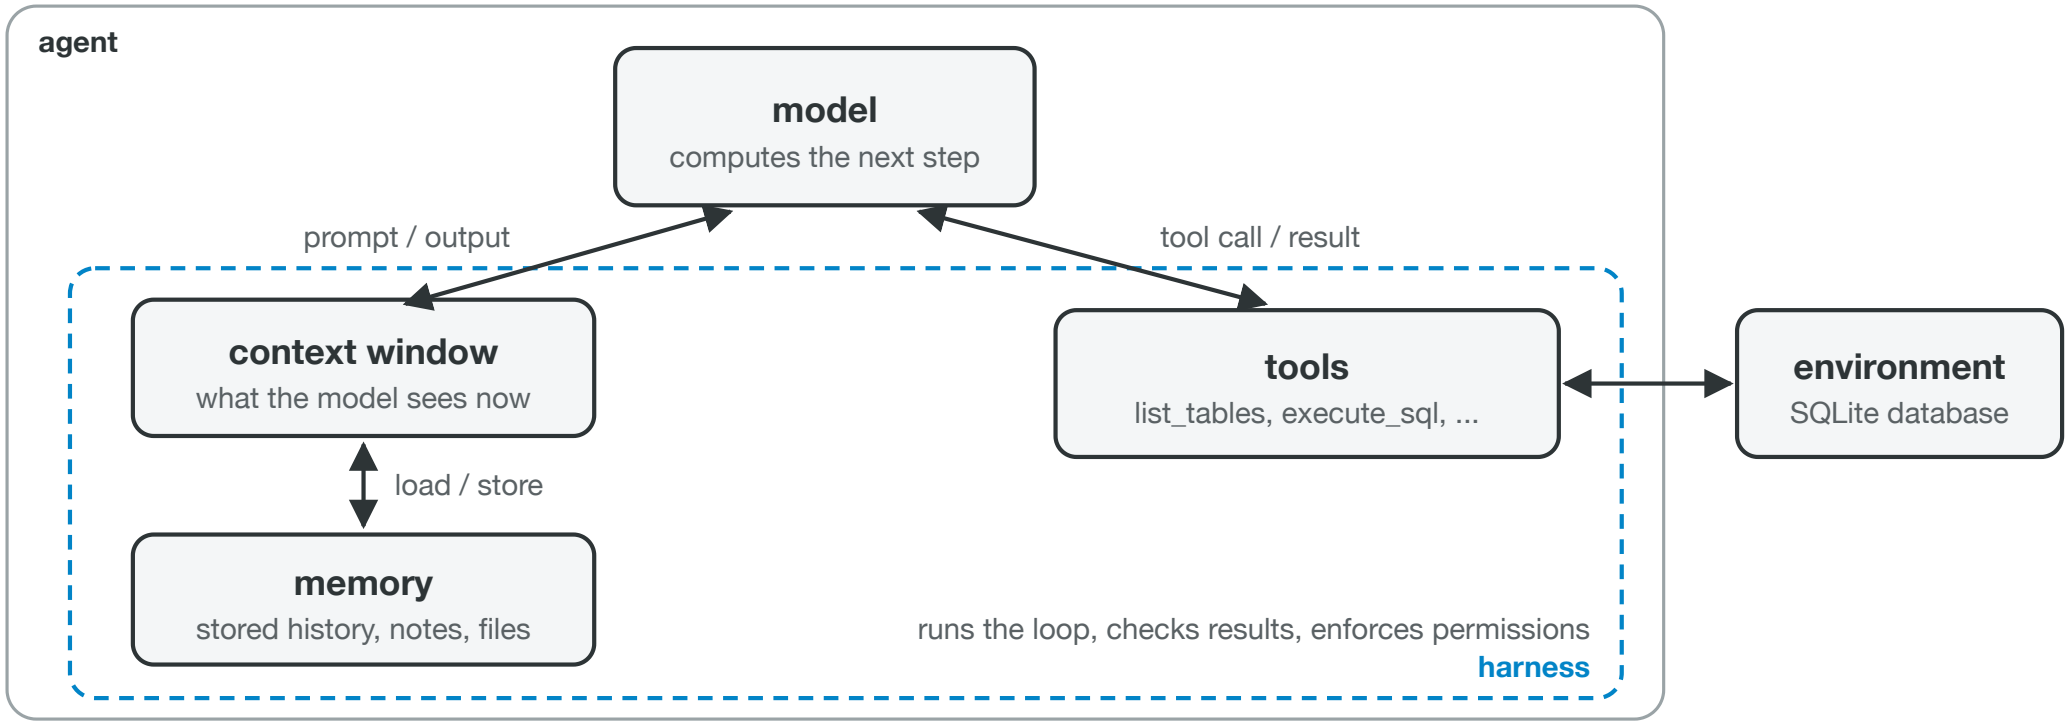
Tel Aviv University

RECAP · HARNESS, LECTURE 1

Foundations of agent harnesses

An agent is a model plus a harness

The model computes; the harness manages its context, memory and tools.



Our SQL agent: 95% of tasks solved vs 60% single-shot, at about 4× the cost.

From Lecture 1 to today

Today we dive deeper and test additional design choices of the harness.

Lecture 1	SQL agent	Today
Execution loop	one short loop	H1 AutoResearch, H2 plan-and-execute, H3 context-managed
Tools	`execute_sql`, and more	`evaluate_equation`, delegation
Context	full history	H3 curated context, fresh executors
Memory	same as context	H3 persistent record
Verification	execution check, LLM judge	training fit, hidden tests, and more
Evaluation	success, cost, and more	ID and OOD fit, and more

PART 1

From components to architectures

Advanced design choices

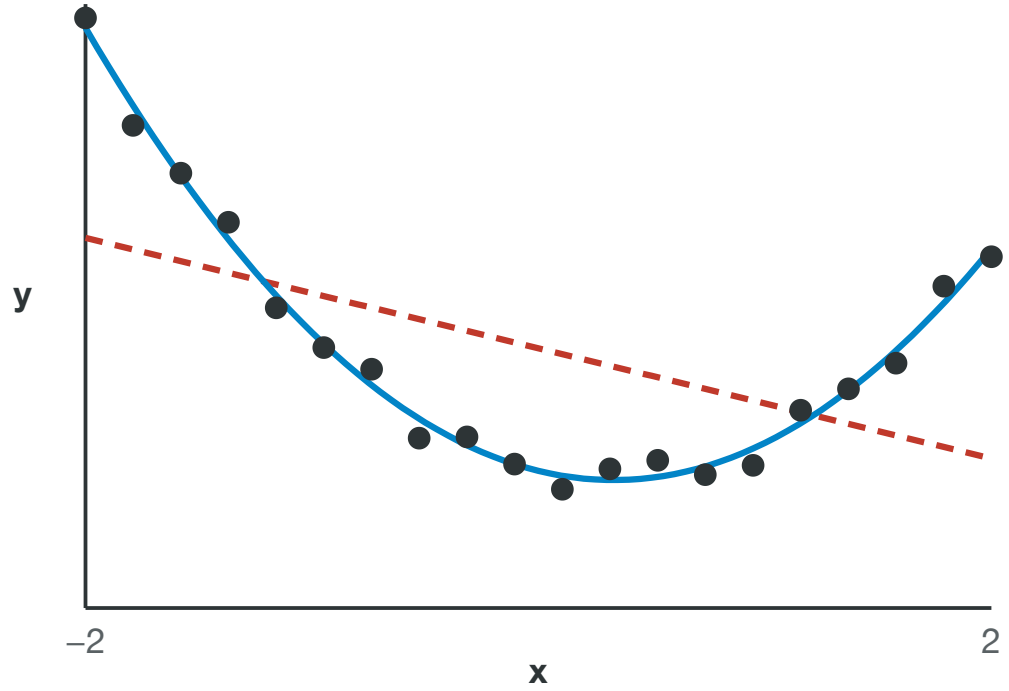
Organizing the components is still an open design decision.

- **Planning:** choose the next action reactively, or plan several investigations?
- **Multi-agent orchestration:** one agent, or several coordinated agents?
- **Context:** keep the full history, or retain only what is relevant?

When does added complexity improve performance enough to justify its cost?

Running example: symbolic regression

Given data $D = \{(x_i, y_i)\}_{i=1}^{n}$, find an expression f such that $y \approx f(x)$.



$f_1(x) = ax + b$

$R^2 = 0.27 \rightarrow$ try another form

$f_2(x) = ax^2 + bx + c$

$R^2 = 0.99$

R^2 returned by evaluate_equation

Shojaee et al., *LLM-SR: Scientific Equation Discovery via Programming with Large Language Models*, ICLR 2025
Shojaee et al., *LLM-SRBench: A New Benchmark for Scientific Equation Discovery with Large Language Models*, ICML 2025

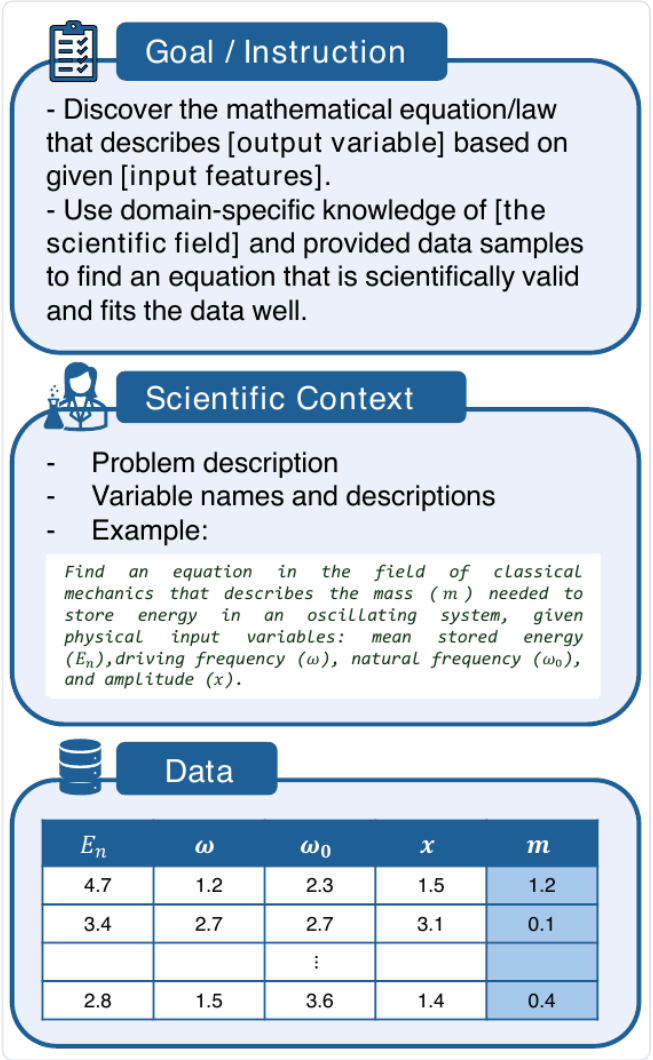
What the agent has

A task description, training data and one tool.

- **Task** q: goal, problem description, variable names
- **Training data** $D_{\text{train}} = \{(x_i, y_i)\}_{i=1}^{n}$
- **Tool:** `evaluate_equation` fits the free parameters θ of a candidate $f(x;\theta)$ and returns its training fit (R^2, NMSE) and the fitted constants

$$\hat{\theta} = \arg\min_{\theta} \sum_{i=1}^{n} \left(y_i - f(x_i;\theta)\right)^2$$

Which other tools may be useful?



The task as an MDP

MDP	Symbolic regression
state s_t	history h_t: the task and all observations so far
action a_t	a candidate equation f_t, sent to `evaluate_equation`
transition	the tool returns o_t: training fit and fitted constants; $h_{t+1} = h_t \cup \{a_t, o_t\}$
proxy reward	$R_{\text{ID}}(\hat{f})$: the fit in the training range, what the agent optimizes
true reward	$R_{\text{OOD}}(\hat{f})$: the fit beyond the training range, what we want
policy π	$\pi_{M,H}$: model M inside harness H

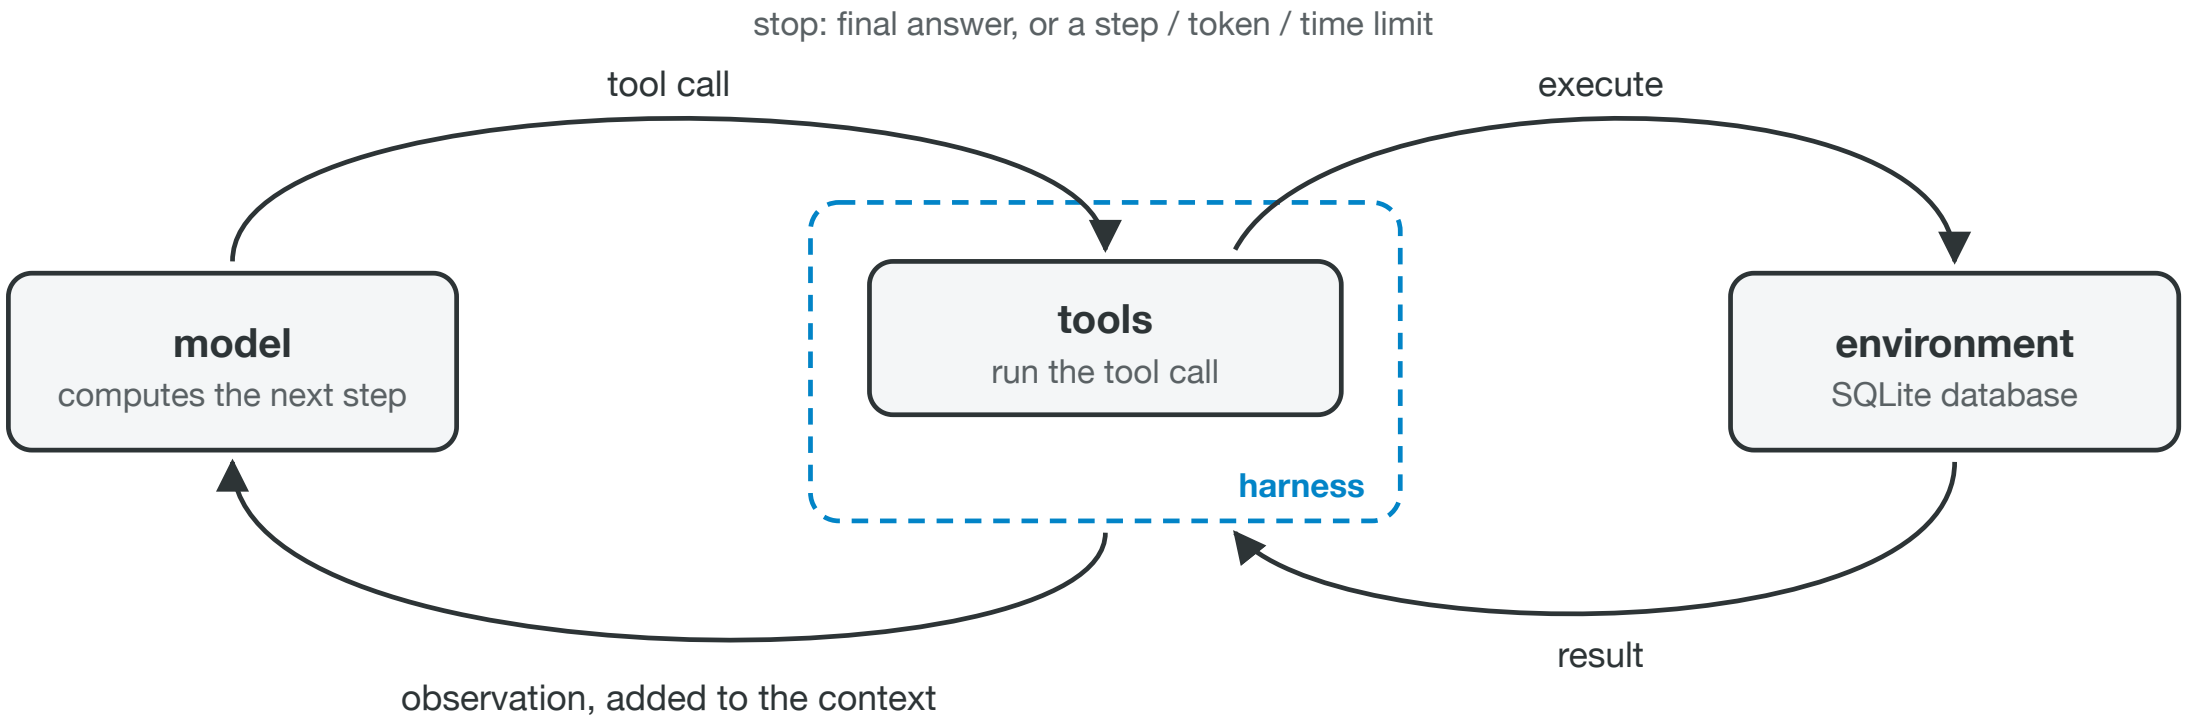
$$J(\pi_{M,H}) = \mathbb{E}\left[R_{\text{ID}}(\hat{f})\right], \qquad J^{\star}(\pi_{M,H}) = \mathbb{E}\left[R_{\text{OOD}}(\hat{f})\right]$$

PART 2

Three harness architectures

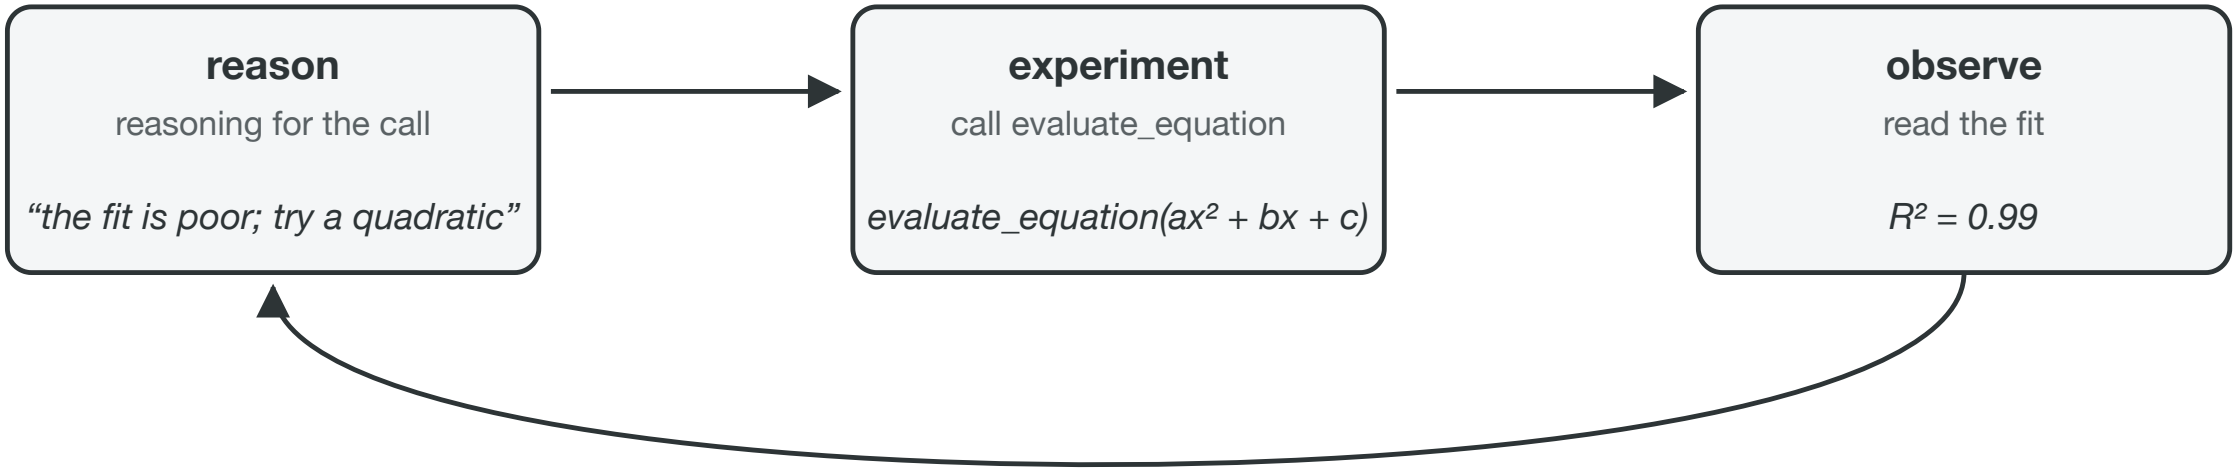
Recall: the execution loop

In Lecture 1, the model proposed actions and the harness executed them.



AutoResearch: the loop, applied to research

AutoResearch runs the same loop, with experiments as actions.



We test AutoResearch both with and without the reasoning step.

Think, then act: the ReAct pattern. Yao et al., *ReAct: Synergizing Reasoning and Acting in Language Models*, ICLR 2023

AutoResearch: the algorithm

One agent picks each next experiment from the full history of experiments.

$$\text{reasoning}_t \sim M(\cdot \mid q, h_t), \qquad f_t \sim M(\cdot \mid q, h_t, \text{reasoning}_t),$$

$$o_t = \texttt{evaluate_equation}(f_t), \quad h_{t+1} = h_t \cup \{\text{reasoning}_t, f_t, o_t\}$$

```
h = [task]
while budget left and not finished:
    reasoning = M(h)              # 1. ask for reasoning
    f = M(h + [reasoning])        # 2. given it, choose the evaluation
    o = evaluate_equation(f)      # fit on training data
    h.append((reasoning, f, o))
return best f in h
```

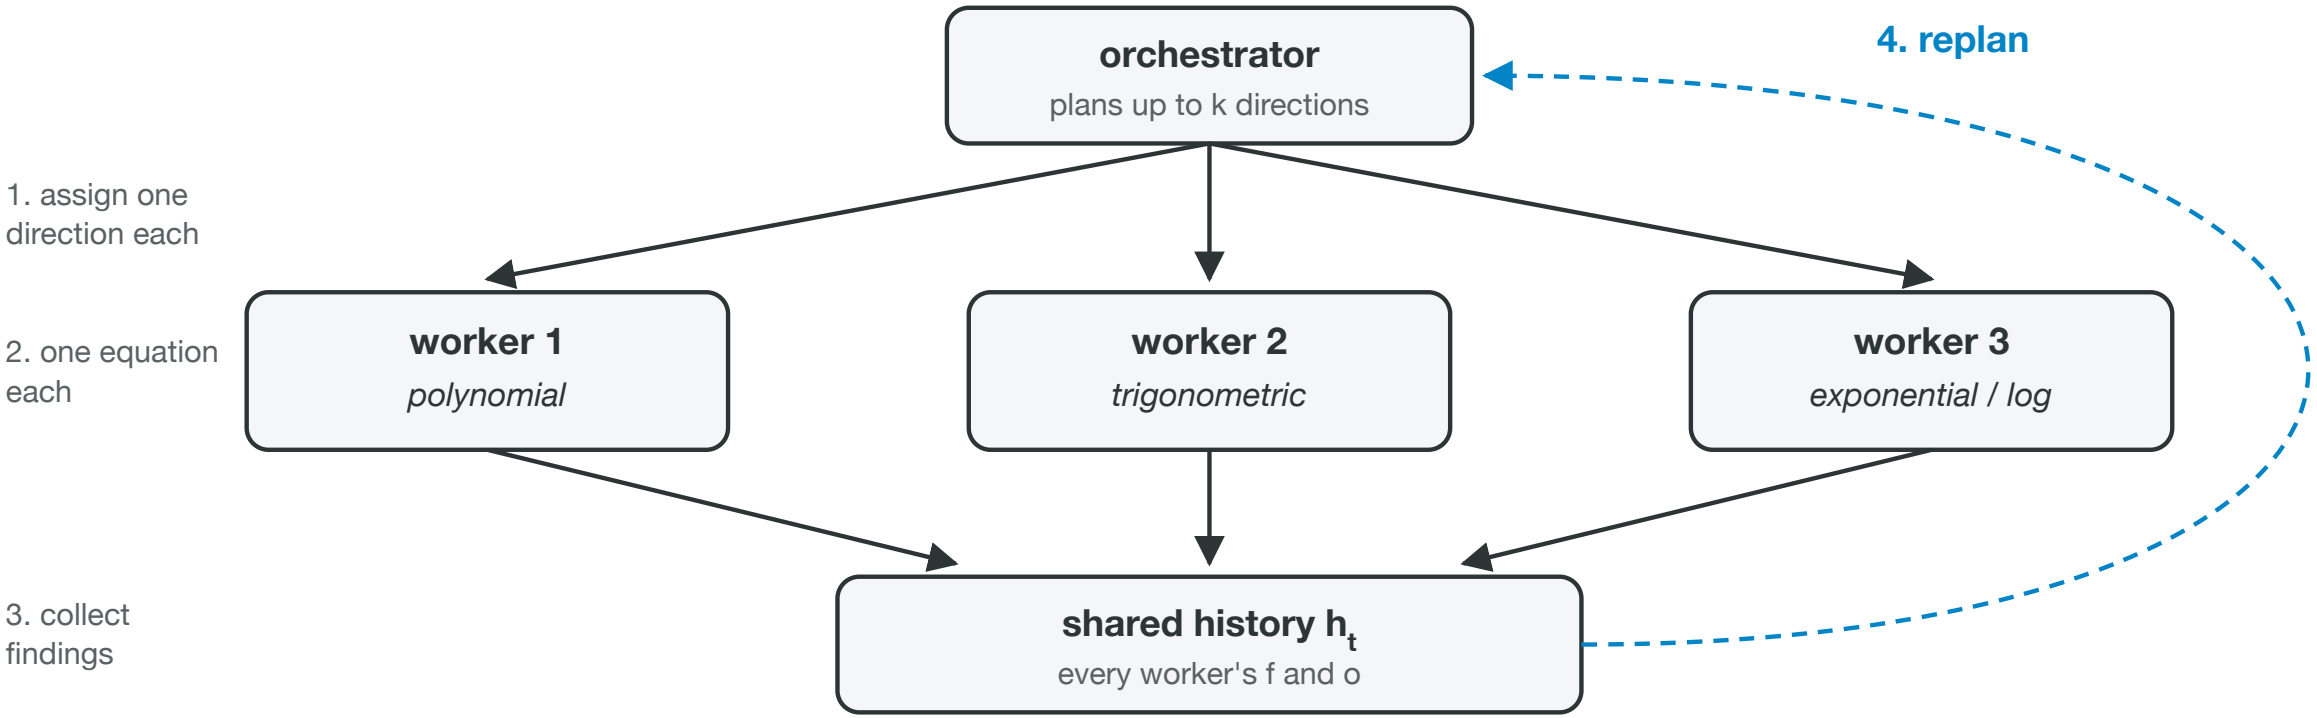
AutoResearch: what we want to learn

- **Reasoning:** how much does it help over a naive baseline with no reasoning?
- **Long runs (test-time compute):** do more experiments keep improving the result?

What else would we like to understand about this design choice?

Plan-and-execute: architecture

An orchestrator plans several research directions and assigns one to each worker.



Plan-and-execute: the algorithm

Each round proposes a set of investigations, runs them, and replans from their results.

$$P_t = \{p_{t,1}, \ldots, p_{t,k}\} \sim M_{\text{orch}}(\cdot \mid q, h_t), \qquad f_{t,j} \sim M(\cdot \mid q, p_{t,j}),$$

$$o_{t,j} = \texttt{evaluate_equation}(f_{t,j}), \qquad h_{t+1} = h_t \cup \{P_t, (f_{t,j}, o_{t,j})_{j=1}^{k}\}$$

```
h = [task]
while budget left and not finished:
    P = M_orch(h)                          # model call: up to k research directions
    F = [M(p) for p in P]                  # one model call per worker: one equation each
    O = [evaluate_equation(f) for f in F]  # fit on training data
    h.append((P, F, O))                    # shared history; replan next round
return best f in h
```

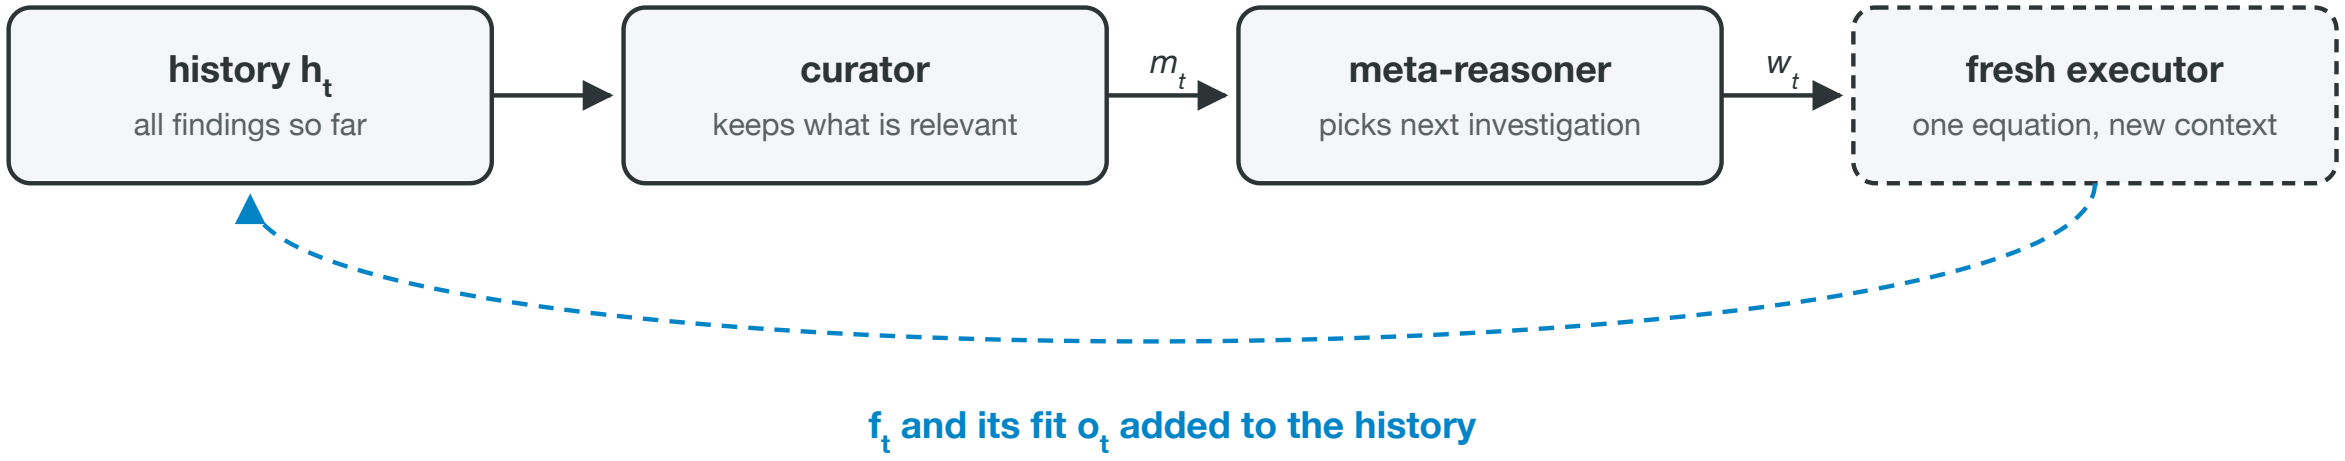
Plan-and-execute: what we want to learn

- **Speedup:** do k workers converge k times faster, in about $1/k$ of the rounds?
- **Task distribution:** does the orchestrator send different workers to explore different options?

What else would we like to understand about this design choice?

Context-managed research: architecture

A meta-reasoner picks the next investigation; a fresh executor runs it.



Context-managed research: the algorithm

Only a curated view of the history reaches each decision.

$$m_t \sim M_{\text{cur}}(\cdot \mid h_t), \quad w_t \sim M_{\text{meta}}(\cdot \mid m_t), \quad f_t \sim M_{\text{exec}}(\cdot \mid q, w_t),$$
$$o_t = \texttt{evaluate_equation}(f_t), \quad h_{t+1} = h_t \cup \{w_t, f_t, o_t\}$$

m_t: the context for this decision, built from the history · w_t: a work order for the executor

```
h = [task]
while budget left and not finished:
    m = M_cur(h)             # model call: build the context for this decision
    w = M_meta(m)            # model call: choose the next investigation
    f = M_exec(w)            # model call, fresh context: one equation
    o = evaluate_equation(f) # fit on training data
    h.append((w, f, o))      # no model call: the history keeps everything
return best f in h
```

In MIRA, each step runs a full investigation: the executor is an agent with many tool calls. Here it proposes a single equation.

MIRA: Learning What to Investigate Next — Meta-Reasoning for Long-Horizon Research Agents, 2026 (arXiv:2610.02525)

Context-managed: what we want to learn

- **Curation vs full history:** does curation beat AutoResearch as runs get longer?

What else would we like to understand about this design choice?

Recap: three architectures

1. **AutoResearch**
2. **Plan-and-Execute**
3. **Context-Managed**

PART 3

Evaluating harness architectures

Which metrics should we measure?

- **ID:** R_{ID} of the final equation $\hat{f}$ on new points in the training range
- **OOD:** R_{OOD} of $\hat{f}$ on points beyond the training range
- **Tokens:** T_H, summed over all model calls

Each is reported against the total budget of each harness.

Which other metrics should we measure?

In distribution and out of distribution

A good fit on the training range does not guarantee a good fit outside it.

- **ID:** R^2 on new points from the training range
- **OOD:** R^2 on points beyond it, e.g. train $x \in [-2, 2]$, OOD $x \in [3, 5]$

Report them separately: an equation can interpolate well and still miss the law.

Cost

Count every model call, not only the ones that test an equation.

$$T_H = \sum_{\text{calls } c} \left(n_c^{\text{in}} + n_c^{\text{out}} \right), \qquad c \in \{\text{executor}, \text{orchestrator}, \text{curator}, \text{meta-reasoner}\}$$

- Also report: API cost, tool calls, equation evaluations

Additional metrics: interpretability

An LLM judge scores what a numeric metric cannot.

- Is the equation scientifically meaningful and understandable?
- Is it consistent with the domain description?

We do not cover this metric here.

PART 4 · CASE STUDY

Three harnesses on LLM-SRBench

LLM-SRBench and the setup

Same model, tool and data; only the harness changes.

Tasks	LLM-SRBench LSR-Synth: chemistry, biology, physics, materials science
Model	`gpt-oss-20b` (open weights) via OpenRouter, for every role
Tool	`evaluate_equation`: least-squares fit of up to 8 constants
Experiment 1	AutoResearch with vs without reasoning: 129 tasks, 5 evaluations
Experiment 2	Plan-and-execute with $k = 1, 2, 4$ workers: 40 tasks, 32 evaluations
Seeds	2 per task and harness

Automatic harness design

$$\max_{f} R(f) \quad (\text{solve a task}) \qquad \text{vs} \qquad \max_{H} \mathbb{E}_{q}\left[R(H, q)\right] \quad (\text{design a harness})$$

- **AFlow:** searches over workflows written as code, using evaluation feedback
- **ADAS:** a meta-agent writes new agent designs as code and keeps the best ones

Zhang et al., *AFlow: Automating Agentic Workflow Generation*, ICLR 2025
Hu, Lu and Clune, *Automated Design of Agentic Systems*, ICLR 2025

Summary

Architectural complexity must be justified empirically.

- **Architectures:** AutoResearch, plan-and-execute and context-managed differ in mechanism and cost
- **Metrics:** ID fit, OOD fit and cost can favor different designs
- **Takeaway:** there are many nuanced design questions, and many ways to evaluate them

With frontier models already tuned to their developers' harnesses, does a custom harness beat a well-written instructions file?